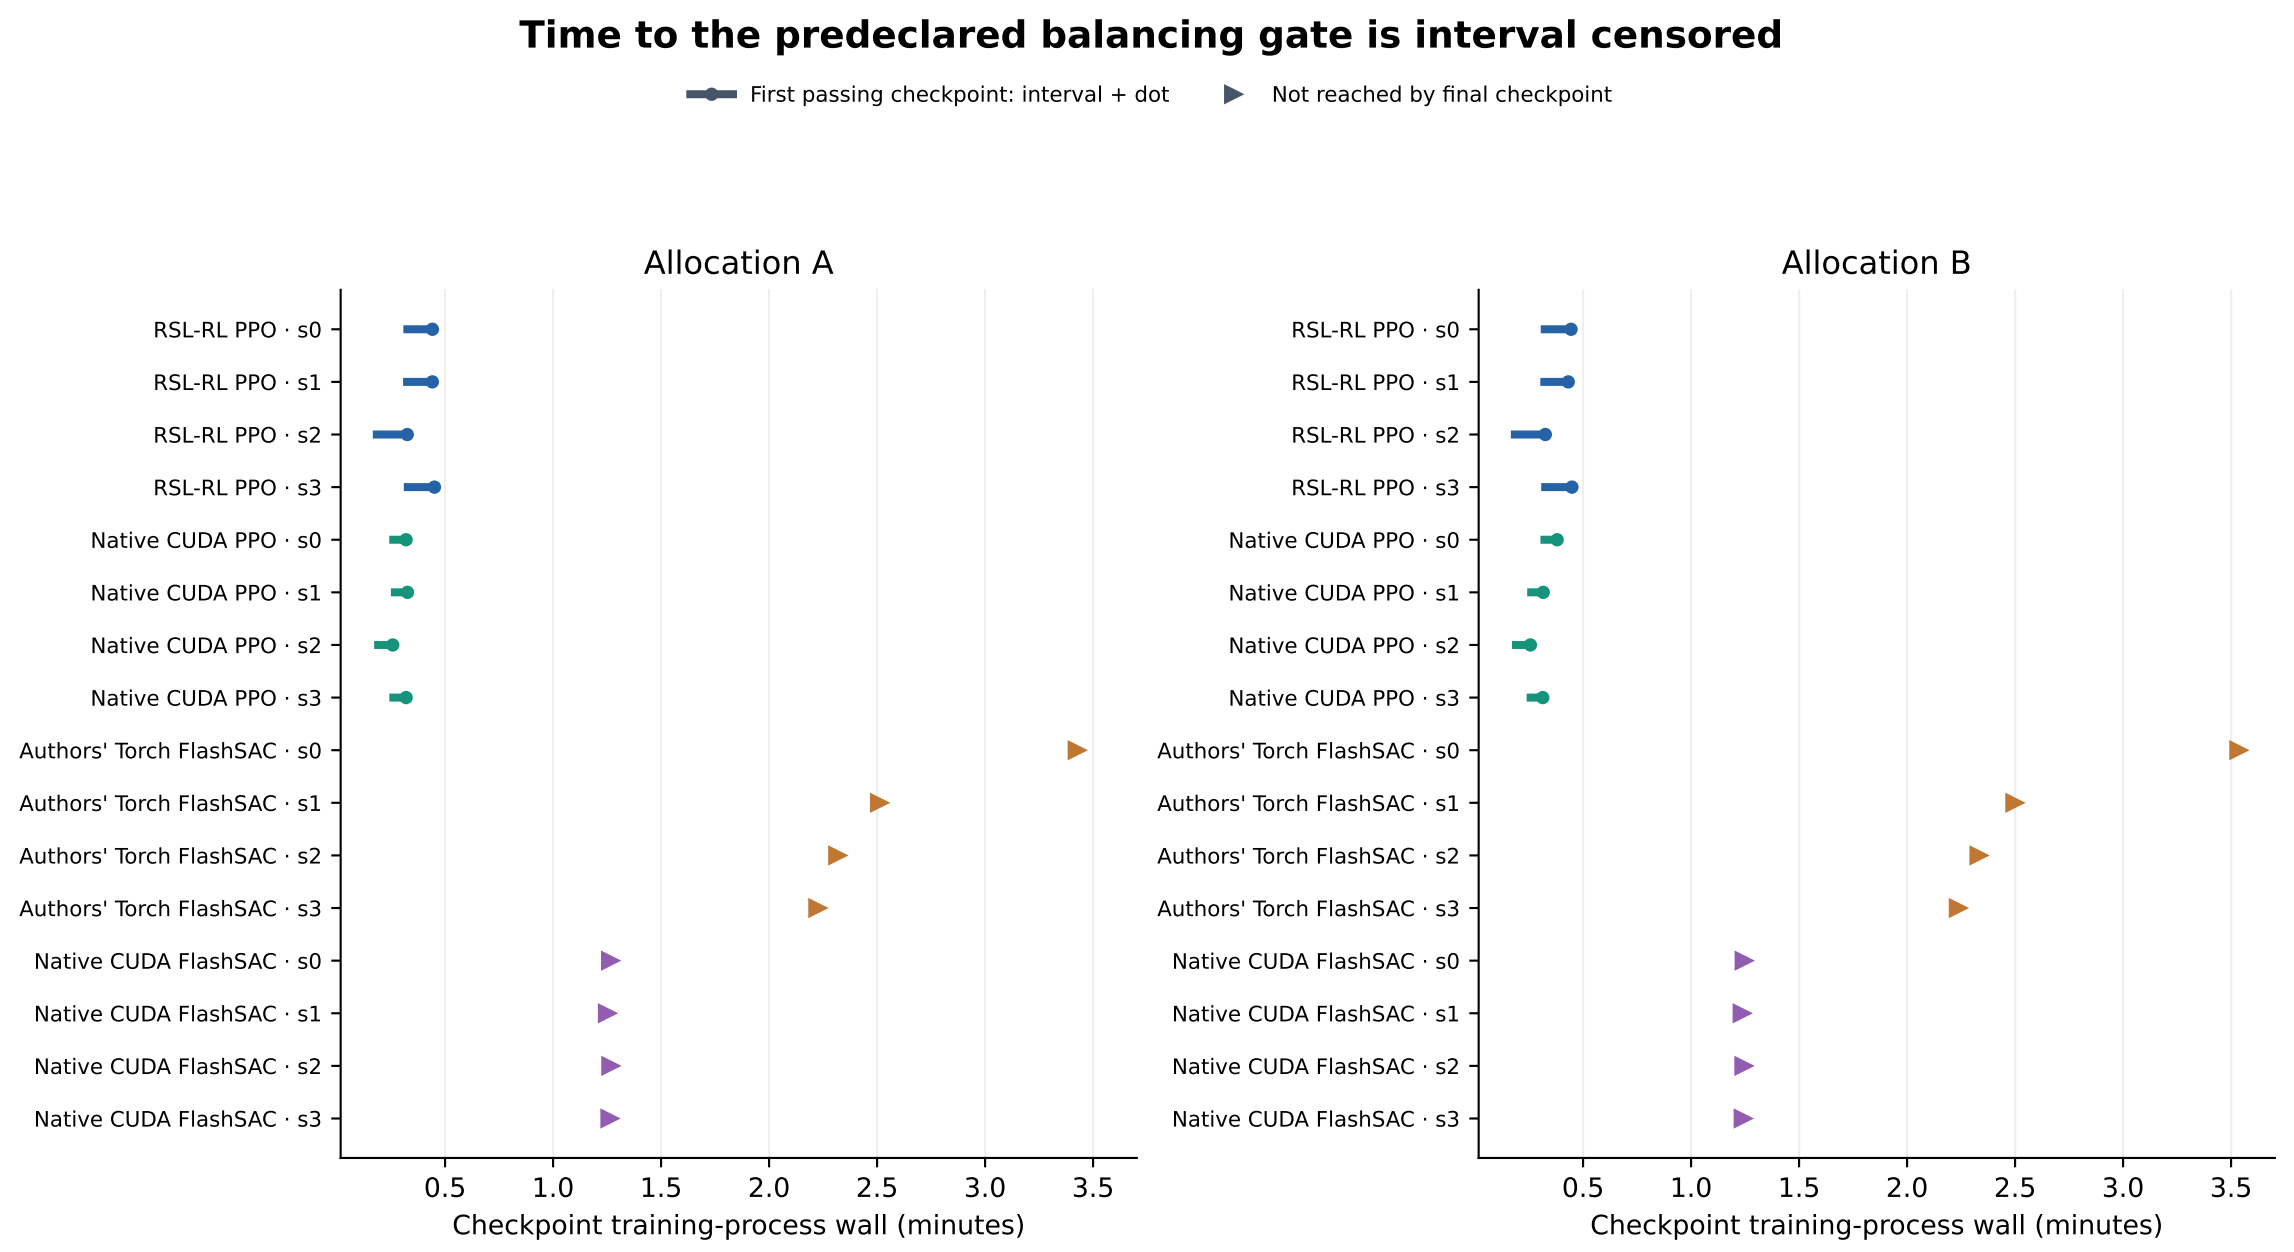
Time to the predeclared balancing gate is interval censored
First passing checkpoint: interval + dot
Not reached by final checkpoint
Allocation A
Allocation B
RSL-RL PPO · s0
RSL-RL PPO · s1
RSL-RL PPO · s2
RSL-RL PPO · s3
Native CUDA PPO · s0
Native CUDA PPO · s1
Native CUDA PPO · s2
Native CUDA PPO · s3
Authors' Torch FlashSAC · s0
Authors' Torch FlashSAC · s1
Authors' Torch FlashSAC · s2
Authors' Torch FlashSAC · s3
Native CUDA FlashSAC · s0
Native CUDA FlashSAC · s1
Native CUDA FlashSAC · s2
Native CUDA FlashSAC · s3
0.5
1.0
1.5
2.0
2.5
3.0
3.5
Checkpoint training-process wall (minutes)
NVIDIA RTX PRO 6000 Blackwell Server Edition · two GPUs (same node) · same four seeds repeated
RSL PPO: FP32, TF32 allowed · Native PPO: strict FP32 · Authors FlashSAC: compiled FP16 AMP · Native FlashSAC: FP16 dense / FP32 state