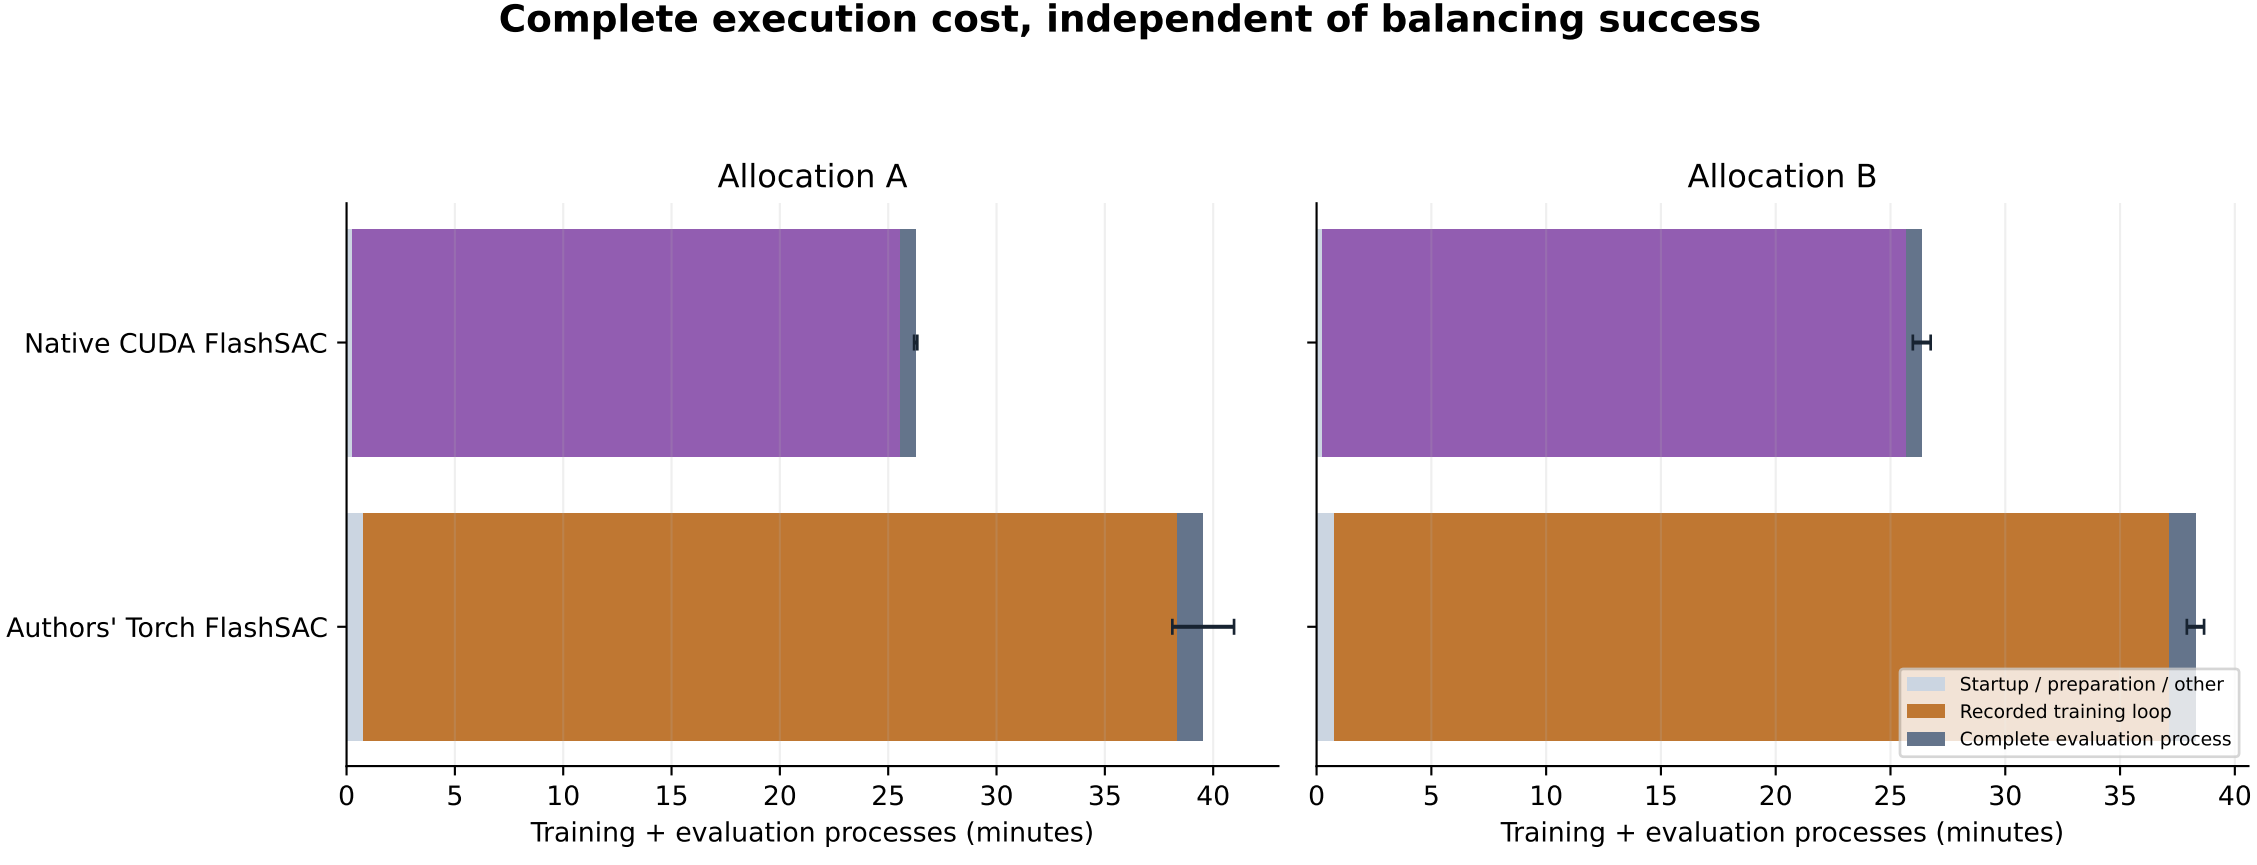
Complete execution cost, independent of balancing success
Allocation A
Allocation B
Native CUDA FlashSAC
Authors' Torch FlashSAC
0
5
10
15
20
25
30
35
40
Training + evaluation processes (minutes)
Startup / preparation / other
Recorded training loop
Complete evaluation process
NVIDIA RTX PRO 6000 Blackwell Server Edition · two distinct GPUs · seeds A0/1, B2/3 (not device replications)
Authors: compiled Torch FP16 AMP · Native: FP16 dense / FP32 state, fused, captured, fresh gradients